

Lab 1: Image Basics and Simple Transformations

(10% of final score, due by Jan. 29 11:59 PM)

In the following tasks, you can use your input images to test your codes. Upload your results and *.py files to the folder of **Lab 1** under **Assignment** in the D2L system. Name the *.py files according to the series numbers of the questions. The *.py files should follow the PEP-8 Style Guide for Python Code [<https://www.python.org/dev/peps/pep-0008/>]

1. Search the Web for job openings in *Computer Vision*. List three jobs that you have interests in, along with the qualifications needed to apply. Summarize in your words instead of copying them here.

2. Read Pillow Image Module: <https://pillow.readthedocs.io/en/stable/reference/Image.html>

Complete the following operations:

- a. Read a RGB image, convert it to a grey level image, and save it as a PNG file.
 - b. Resize the grey level image to half width and half height and show the resized image on screen.
3. Plotting Numpy arrays as images. Set up a 4×5 (rows by columns) array and plot the array as a gray-level 8-bit image, using matplotlib.pyplot
[https://matplotlib.org/3.5.3/api/_as_gen/matplotlib.pyplot.html]
 4. Suppose an image has 640 columns and 480 rows and is stored in row-major order. Convert the coordinates $(x, y) = (38, 52), (592, 241),$ and $(33, 0)$ to 1D indices. Conversely, convert the following 1D indices to (x, y) coordinates: $i = 8092, 24061,$ and 38190 .

5. Given the following 8-bit grayscale image,

$$\begin{bmatrix} 176 & 94 & 201 & 219 \\ 37 & 161 & 16 & 88 \\ 71 & 129 & 177 & 81 \\ 41 & 198 & 107 & 19 \end{bmatrix}$$

- a. Perform the following operations **by hand**: flip, flop, invert, and rotate clockwise by 90 degrees. Write down the new arrays.
- b. Follow the pseudocode provided below to implement the same operations in Python for a grayscale image: flip, flop, invert, and rotate clockwise by 90 degrees.

Here is the pseudocode for reference:

ALGORITHM 3.1 Flip an image by reflecting about a horizontal axis

FLIPIMAGE (I)

Input: image I of size $width \times height$

Output: upside-down image I'

```

1   $I' \leftarrow \text{ALLOCATEIMAGE}(width, height)$ 
2  for  $(x, y) \in I$  do
3       $I'(x, height - 1 - y) \leftarrow I(x, y)$ 
4  return  $I'$ 

```

- Allocate memory for output image.
- For each pixel in input image, set corresponding pixel in output image.
- Return output image.

ALGORITHM 3.2 Flop an image by reflecting about a vertical axis

FLOPIMAGE (I)

Input: image I of size $width \times height$

Output: mirror-reversed image I'

```

1   $I' \leftarrow \text{ALLOCATEIMAGE}(width, height)$ 
2  for  $(x, y) \in I$  do
3       $I'(width - 1 - x, y) \leftarrow I(x, y)$ 
4  return  $I'$ 

```

- Allocate memory for output image.
- For each pixel in input image, set corresponding pixel in output image.
- Return output image.

ALGORITHM 3.3 Rotate an image by a multiple of 90 degrees

ROTATEIMAGEBYMULTIPLEOF90DEGREES(I, m)

Input: image I of size $width \times height$, signed integer m indicating the number of 90-degree turns

Output: image I' of size $new-width \times new-height$, which is I rotated by $90m$ degrees clockwise

```

1  case  $\text{mod}(m, 2)$  of
2      0:  $new-width \leftarrow width, new-height \leftarrow height$ 
3      1:  $new-width \leftarrow height, new-height \leftarrow width$ 
4   $I' \leftarrow \text{ALLOCATEIMAGE}(new-width, new-height)$ 
5  for  $(x', y') \in I'$  do
6      case  $\text{mod}(m, 4)$  of
7          0:  $I'(x', y') \leftarrow I(x', y')$ 
8          1:  $I'(x', y') \leftarrow I(y', height - 1 - x')$ 
9          2:  $I'(x', y') \leftarrow I(width - 1 - x', height - 1 - y')$ 
10         3:  $I'(x', y') \leftarrow I(width - 1 - y', x')$ 
11 return  $I'$ 


```

- Image dimensions remain the same.
- Image dimensions are swapped.
- no rotation
- 90 degrees clockwise
- 180 degrees
- 90 degrees counterclockwise

- c. Modify the algorithm to implement the same operations for a RGB color image, and plot the four outputs (the flip one, the flop one, the invert one, and the rotate one) in a 2x2 subplot array, using Python **Matplotlib** Module.

[https://matplotlib.org/tutorials/introductory/sample_plots.html]

The output would look like this (you can use your own color image):



6. Given the following 4-bit image:

$$\begin{bmatrix} 5 & 8 & 3 & 7 \\ 1 & 3 & 3 & 9 \\ 6 & 8 & 2 & 7 \\ 4 & 1 & 0 & 9 \end{bmatrix}$$

- Compute the histogram, normalized histogram, and cumulative normalized histogram for the image by hand and provide calculation process.
- Follow the pseudocode provided below to implement histogram equalization in Python and run the code on a grayscale image.

Here is the pseudocode for reference:

ALGORITHM 3.9 Perform histogram equalization on an image

HISTOGRAMEQUALIZE(I)

Input: grayscale image I

Output: histogram-equalized grayscale image I' with increased contrast

```

1  $\bar{h} \leftarrow \text{COMPUTENORMALIZEDHISTOGRAM}(I)$ 
2  $\bar{c} \leftarrow \text{RUNNINGSUM}(\bar{h})$ 
3 for  $(x, y) \in I$  do
4    $I'(x, y) \leftarrow \text{ROUND}(255 * \bar{c}[I(x, y)])$ 
5 return  $I'$ 

```

RUNNINGSUM(a)

Input: 1D array a of $length$ values

Output: 1D running sum s of array

```

1  $s[0] \leftarrow a[0]$ 
2 for  $k \leftarrow 1$  to  $length - 1$  do
3    $s[k] \leftarrow s[k - 1] + a[k]$ 
4 return  $s$ 

```

- c. Search Python scikit-image module or openCV module to find a function that implements histogram equalization. Compare your result with their results pixel-by-pixel. Output the comparison in a black-and-white image: black means no difference while white means different.

7. Given the following 3 consecutive frames,

$$I1 = \begin{bmatrix} 18 & 168 & 94 & 67 \\ 120 & 97 & 78 & 198 \\ 83 & 70 & 208 & 17 \\ 238 & 208 & 189 & 68 \end{bmatrix} \quad I2 = \begin{bmatrix} 21 & 168 & 92 & 71 \\ 122 & 71 & 191 & 227 \\ 83 & 212 & 16 & 187 \\ 240 & 216 & 188 & 68 \end{bmatrix} \quad I3 = \begin{bmatrix} 20 & 171 & 92 & 70 \\ 76 & 193 & 39 & 255 \\ 209 & 20 & 20 & 194 \\ 241 & 210 & 190 & 73 \end{bmatrix}$$

- a. Compute the double difference and the triple difference between the 3 frames by hand, using the threshold $\tau = 40$;
- b. Follow the pseudocode provided below to implement the double difference and the triple difference in Python and run the code on three consecutive frames to detect a moving object. Here are the pseudocodes for reference.

ALGORITHM 3.13 Compute the double difference between three consecutive image frames

FRAMEDIFFERENCEDOUBLE ($I_{t-1}, I_t, I_{t+1}, \tau$)

Input: successive images I_{t-1}, I_t , and I_{t+1} , and threshold τ

Output: binary image indicating the moving regions

```

1  for  $(x, y) \in I_t$  do
2       $d_1 \leftarrow |I_{t-1}(x, y) - I_t(x, y)|$ 
3       $d_2 \leftarrow |I_{t+1}(x, y) - I_t(x, y)|$ 
4       $I'(x, y) \leftarrow 1$  if  $d_1 > \tau$  AND  $d_2 > \tau$  else 0
5  return  $I'$ 
```

ALGORITHM 3.14 Compute the triple difference between three consecutive image frames

FRAMEDIFFERENCETRIPLE ($I_{t-1}, I_t, I_{t+1}, \tau$)

Input: successive images I_{t-1}, I_t , and I_{t+1} , and threshold τ

Output: binary image indicating the moving regions

```

1  for  $(x, y) \in I_t$  do
2       $d_1 \leftarrow |I_{t-1}(x, y) - I_t(x, y)|$ 
3       $d_2 \leftarrow |I_{t+1}(x, y) - I_t(x, y)|$ 
4       $d_3 \leftarrow |I_{t+1}(x, y) - I_{t-1}(x, y)|$ 
5       $I'(x, y) \leftarrow 1$  if  $d_1 + d_2 - d_3 > \tau$  else 0
6  return  $I'$ 
```

8. Given two images with masks:

$$I_A = \begin{bmatrix} 132 & 231 & 227 \\ 237 & 105 & 238 \\ 193 & 59 & 128 \end{bmatrix}, \quad M_A = \begin{bmatrix} 255 & 255 & 0 \\ 255 & 255 & 0 \\ 0 & 0 & 0 \end{bmatrix}$$

$$I_B = \begin{bmatrix} 43 & 79 & 116 \\ 56 & 246 & 184 \\ 36 & 119 & 162 \end{bmatrix}, \quad M_B = \begin{bmatrix} 0 & 0 & 0 \\ 0 & 255 & 255 \\ 0 & 255 & 255 \end{bmatrix}$$

Compute by hand with the three Porter-Duff operators: I_A OVER I_B , I_A IN I_B , and I_A ATOP I_B . Show only the values of the valid pixels; indicate invalid pixels using an X.

9. Given the following image

$$I = \begin{bmatrix} 232 & 177 & 82 & 7 \\ 241 & 18 & 152 & 140 \\ 156 & 221 & 67 & 3 \end{bmatrix}$$

- Use bilinear interpolation to compute the value at $(0.1, 0.7)$, $(1.2, 0.5)$, $(1.3, 1.6)$, and $(2.8, 1.7)$ **by hand** and provide the calculation process.
- Follow the pseudocode provided below to implement the bilinear interpolation in Python and run the code on the given 4×3 image.

Here is the pseudocode for reference:

ALGORITHM 3.16 Perform bilinear interpolation on an image at a point

INTERPOLATEBILINEAR(I, x, y)

Input: image, floating-point coordinates (x, y)

Output: weighted average of gray levels of nearest four pixels

```
1   $x_0 \leftarrow \text{FLOOR}(x)$ 
2   $y_0 \leftarrow \text{FLOOR}(y)$ 
3   $\alpha_x \leftarrow x - x_0$ 
4   $\alpha_y \leftarrow y - y_0$ 
5   $\bar{\alpha}_x \leftarrow 1 - \alpha_x$ 
6   $\bar{\alpha}_y \leftarrow 1 - \alpha_y$ 
7  return  $\bar{\alpha}_x * \bar{\alpha}_y * I(x_0, y_0) + \alpha_x * \bar{\alpha}_y * I(x_0 + 1, y_0) + \bar{\alpha}_x * \alpha_y * I(x_0, y_0 + 1) + \alpha_x * \alpha_y * I(x_0 + 1, y_0 + 1)$ 
```