

Lab 6: Object Detection with Deep NN

(10% of final score, due by ~~Apr 8~~, 11:59 PM)

Apr. 9

Lab 6 will be implemented on Google CoLab. Please keep the code file output when saving the notebook and share the code with: abbas.salehitangrizi@gmail.com. Please add comments if needed and follow the PEP-8 Style Guide for Python Code. [<https://www.python.org/dev/peps/pep-0008/>]

1. Convolutional Neural Network (CNN): <https://www.tensorflow.org/tutorials/images/cnn>

Read the example file <cnn.ipynb> , save it to your G-drive, and run it in Google CoLab. Answer the following questions:

- How is the sequential model defined?
- Read the documentation of `tf.keras.layers.Conv2D`.
https://www.tensorflow.org/api_docs/python/tf/keras/layers/Conv2D
If the inputs are 50×50 RGB images, `y= tf.keras.layers.Conv2D()`, what would the `y.shape` be with: batch size = 4, strides = 2 , and padding="valid"
- Before adding two dense layers, how is the number of total parameters 56,320 calculated?
- After the sequential model, why are two dense layers added?
- After the two dense layers, the number of total parameters is 122,570. How is it calculated?
- This CNN achieved a test accuracy of about 70%. Is it a good number? How to improve it?

2. Object Detection Models: https://www.tensorflow.org/hub/tutorials/tf2_object_detection

Read the example file <Object Detection Inference on TF 2 and TF Hub.ipynb>, save it to your G-drive, and run it in Google Colab.

- Compare models: “Mask R-CNN Inception ResNet V2 1024×1024” and “EfficientDet D0 521×512.” Evaluation metric: inference time. Using the same test image. Describe your observation and explain it.
- Do some research on the models “Mask R-CNN Inception ResNet V2 1024×1024” and “EfficientDet D7 1536×1536” and briefly (one paragraph for each) describe the mechanism of the models. Why is EfficientDet very efficient? Why does Mask R-CNN allow instance segmentation? Here is the link to the original publications:
https://openaccess.thecvf.com/content_CVPR_2020/papers/Tan_EfficientDet_Scalable_and_Efficient_Object_Detection_CVPR_2020_paper.pdf
https://openaccess.thecvf.com/content_ICCV_2017/papers/He_Mask_R-CNN_ICCV_2017_paper.pdf

3. Transfer learning:

https://colab.research.google.com/github/tensorflow/hub/blob/master/examples/colab/tf2_image_retraining.ipynb

Read the example file <TF Hub for TF2: Retraining an image classifier.ipynb>, save it to your G-drive, and run it in Google Colab.

- a) `model.summary()` outputs some parameter are trainable while others are non-trainable. Why? Check the box “do_fine_tuning” and its value will be changed to True. Run the model again. The outputs of `model.summary()` will change and the number of trainable parameters will significantly increase. Why?
- b) The default `batch_size` is 16. The parameters `steps_per_epoch` and `validation_steps` vary with `batch_size`. Adjust `batch_size` to 8, 4, and 2, observe inference time, loss, and accuracy, and describe the changes and explain why?
- c) Try out the model on your input images with and without flowers (not from the validation data). Are the predictions correct?